

Hướng dẫn cài đặt ceph 18.2.2

Yêu cầu

- Python 3
- Đã cài đặt Docker để chạy container
- Đồng bộ hóa thời gian (như Chrony hoặc phiên bản cũ `ntpd`)
- LVM2 để cung cấp lưu trữ
- 2 Interface mạng: 1 cho public network, 1 cho cluster network

Chú ý: Hướng dẫn này giả định rằng server đã được cài đặt Chrony và Docker

LAB

3 server chạy Ubuntu 22.04

Interface:

- Public: 192.168.3.0/24
- Cluster: 10.10.10.0/24

ceph-01: public: 192.168.3.71/24; cluster: 10.10.10.71/24; (mon, mgr, osd)

ceph-02: public: 192.168.3.72/24; cluster: 10.10.10.72/24; (mon, mgr, osd)

ceph-03: public: 192.168.3.73/24; cluster: 10.10.10.73/24; (mon, mgr, osd)

Cài đặt cephadm

Trên 3 host xóa hoặc comment dòng 127.0.1.1 trong file **/etc/hosts**

Cephadm sử dụng kết nối SSH để tương tác và triển khai dịch vụ. Cluster được xây dựng bằng cách download tệp binary và chạy bootstrap. Cephadm là cách dễ nhất để triển khai và chạy một cụm Ceph mới một cách nhanh chóng.

Chú ý: Cụm lưu trữ DC Funtap sử dụng Cephadm để triển khai.

Trên node ceph-01:

```
root@ceph-01:~# CEPH_RELEASE=18.2.2
```

```
root@ceph-01:~# curl --silent --remote-name --location https://download.ceph.com/rpm-${CEPH_RELEASE}/el9/noarch/cephadm
```

```
root@ceph-01:~# chmod +x cephadm
```

```
root@ceph-01:~# ./cephadm add-repo --release 18.2.2
```

```
root@ceph-01:~# ./cephadm install
```

```
root@ceph-01:~# CEPH_RELEASE=18.2.2
root@ceph-01:~# curl --silent --remote-name --location https://download.ceph.com/rpm-${CEPH_RELEASE}/el9/noarch/cephadm
root@ceph-01:~# chmod +x cephadm
root@ceph-01:~# ./cephadm add-repo --release 18.2.2
Installing repo GPG key from https://download.ceph.com/keys/release.gpg...
Installing repo file at /etc/apt/sources.list.d/ceph.list...
Updating package list...
Completed adding repo.
root@ceph-01:~# ./cephadm install
Installing packages ['cephadm']...
root@ceph-01:~# which cephadm
/usr/sbin/cephadm
```

```
root@ceph-01:~# docker pull quay.io/ceph/ceph:v18.2.2
```

Triển khai cụm ceph

```
root@ceph-01:~# cephadm --docker --image quay.io/ceph/ceph:v18.2.2
bootstrap --mon-ip 192.168.3.71 --cluster-network 10.10.10.0/24
```

Kết thúc lệnh ceph sẽ trả về user/pass admin của trang quản lý.

```

Adding key to root@localhost authorized_keys...
Adding host ceph-01...
Deploying mon service with default placement...
Deploying mgr service with default placement...
Deploying crash service with default placement...
Deploying ceph-exporter service with default placement...
Deploying prometheus service with default placement...
Deploying grafana service with default placement...
Deploying node-exporter service with default placement...
Deploying alertmanager service with default placement...
Enabling the dashboard module...
Waiting for the mgr to restart...
Waiting for mgr epoch 9...
mgr epoch 9 is available
Generating a dashboard self-signed certificate...
Creating initial admin user...
Fetching dashboard port number...
Ceph Dashboard is now available at:

    URL: https://ceph-01:8443/
    User: admin
    Password: u5oe80lnzo

Enabling client.admin keyring and conf on hosts with "admin" label
Saving cluster configuration to /var/lib/ceph/41073fa0-795e-11ef-bedb-fd1565f7fbc9/config directory
Enabling autotune for osd_memory_target
You can access the Ceph CLI as following in case of multi-cluster or non-default config:

    sudo /usr/sbin/cephadm shell --fsid 41073fa0-795e-11ef-bedb-fd1565f7fbc9 -c /etc/ceph/ceph.conf -k /etc/ceph/ceph.client.admin.keyring

Or, if you are only running a single cluster on this host:

    sudo /usr/sbin/cephadm shell

Please consider enabling telemetry to help improve Ceph:
    ceph telemetry on

For more information see:
    https://docs.ceph.com/en/latest/mgr/telemetry/

Bootstrap complete.

```

```

root@ceph-01:~# docker ps -a
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS                 NAMES
d976784a7d7e        quay.io/prometheus/alertmanager:v0.25.0   "/bin/alertmanager -"   15 minutes ago     Up 15 minutes          ceph-41073fa0-795e-11ef-bedb-fd1565f7fbc9-alertmanager-ceph-01
7ceb5c8a62a9        quay.io/ceph/ceph-grafana:9.4.7      "/bin/sh -c 'grafana'"  15 minutes ago     Up 15 minutes          ceph-41073fa0-795e-11ef-bedb-fd1565f7fbc9-grafana-ceph-01
e4ae52fd673b        quay.io/prometheus/prometheus:v2.43.0     "/bin/prometheus --c"  15 minutes ago     Up 15 minutes          ceph-41073fa0-795e-11ef-bedb-fd1565f7fbc9-prometheus-ceph-01
1b4c1f534088        quay.io/prometheus/node-exporter:v1.5.0   "/bin/node_exporter "  18 minutes ago     Up 18 minutes          ceph-41073fa0-795e-11ef-bedb-fd1565f7fbc9-node-exporter-ceph-01
3ab635abe0f5        quay.io/ceph/ceph                   "/usr/bin/ceph-crash"  18 minutes ago     Up 18 minutes          ceph-41073fa0-795e-11ef-bedb-fd1565f7fbc9-crash-ceph-01
bbe1159cal1c12      quay.io/ceph/ceph                   "/usr/bin/ceph-expor"  18 minutes ago     Up 18 minutes          ceph-41073fa0-795e-11ef-bedb-fd1565f7fbc9-ceph-exporter-ceph-01
da0038f156c1        quay.io/ceph/ceph:v18.2.2       "/usr/bin/ceph-mgr -"   19 minutes ago     Up 19 minutes          ceph-41073fa0-795e-11ef-bedb-fd1565f7fbc9-mgr-ceph-01-iexdoz
9e3d1c58edd7        quay.io/ceph/ceph:v18.2.2       "/usr/bin/ceph-mon -"   19 minutes ago     Up 19 minutes          ceph-41073fa0-795e-11ef-bedb-fd1565f7fbc9-mon-ceph-01

```

Ceph CLI

Có 2 cách để sử dụng CLI:

Cách 1: cài đặt ceph-common

```
root@ceph-01:~# cephadm install ceph-common
```

Cách 2: sử dụng cephadm shell

```
root@ceph-01:~# cephadm shell
```

Lệnh này sẽ tạo và chạy một container để thực thi các lệnh ceph bên trong container.

Xác nhận đã cài thành công:

```
root@ceph-01:~# ceph -v
root@ceph-01:~# ceph -v
ceph version 18.2.2 (531c0d11a1c5d39fbfe6aa8a521f023abf3bf3e2) reef (stable)
root@ceph-01:~#
```

Thêm node mới

```
root@ceph-01:~# ceph cephadm get-pub-key
```

Lệnh này trả về public key. Copy key này vào các node sẽ tham gia vào cluster

```
root@ceph-01:~# ssh-copy-id -f -i /etc/ceph/ceph.pub
root@192.168.3.72
```

```
root@ceph-01:~# ssh-copy-id -f -i /etc/ceph/ceph.pub
root@192.168.3.73
```

```
root@ceph-01:~# ssh-copy-id -f -i /etc/ceph/ceph.pub root@192.168.3.72
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/etc/ceph/ceph.pub"
The authenticity of host '192.168.3.72 (192.168.3.72)' can't be established.
ED25519 key fingerprint is SHA256:sDEujzsFBeAwPXRjLdTKJ5reTMBJ9L3djJiFHnbcWZA.
This host key is known by the following other names/addresses:
  ~/.ssh/known_hosts:18: [hashed name]
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Please type 'yes', 'no' or the fingerprint: yes

Number of key(s) added: 1

Now try logging into the machine, with:  "ssh 'root@192.168.3.72'"
and check to make sure that only the key(s) you wanted were added.

root@ceph-01:~# ssh-copy-id -f -i /etc/ceph/ceph.pub root@192.168.3.73
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/etc/ceph/ceph.pub"
The authenticity of host '192.168.3.73 (192.168.3.73)' can't be established.
ED25519 key fingerprint is SHA256:LWlhSxtCVqymQxzC9550XJQKL07PRdKBErwQFGF6z9I.
This host key is known by the following other names/addresses:
  ~/.ssh/known_hosts:21: [hashed name]
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes

Number of key(s) added: 1

Now try logging into the machine, with:  "ssh 'root@192.168.3.73'"
and check to make sure that only the key(s) you wanted were added.
```

```
root@ceph-01:~# ceph orch host add ceph-02 192.168.3.72 --labels
admin
```

```
root@ceph-01:~# ceph orch host add ceph-03 192.168.3.73 --labels
admin
```

```
root@ceph-01:~# ceph orch host add ceph-02 192.168.3.72 --labels _admin
Added host 'ceph-02' with addr '192.168.3.72'
root@ceph-01:~# ceph orch host add ceph-03 192.168.3.73 --labels _admin
Added host 'ceph-03' with addr '192.168.3.73'
```

Lệnh này sẽ thêm node mới vào cluster đồng thời gắn labels _admin vào node.

labels _admin: để chỉ ra rằng từ các node này có thể chạy lệnh ceph cli hoặc triển khai các daemons dựa trên labels chỉ định.

Chú ý: Tên host phải đúng với hostname của server

Sau khi thêm node hãy chờ 1 lúc để các thành phần như mon, mgr được triển khai trên node mới.

Thêm OSD

OSD chịu trách nhiệm lưu trữ dữ liệu, xử lý các yêu cầu đọc/ghi từ client, và nhân bản dữ liệu để đảm bảo tính toàn vẹn và khả năng chịu lỗi trong cụm Ceph.

Có thể kiểm tra trên các node có ổ cứng nào phù hợp để triển khai OSD không bằng lệnh:

```
root@ceph-01:~# ceph orch device ls
```

Nếu có ổ cứng nào không phù hợp sẽ thấy thông báo lý do ở cột reject

```
root@ceph-01:~# ceph orch device ls
HOST      PATH    TYPE   DEVICE ID                               SIZE  AVAILABLE  REFRESHED  REJECT  REASONS
ceph-01   /dev/sda hdd    ATA_ST1000DM010-2EP1_ZN14H4AG    931G Yes        13m ago
ceph-02   /dev/sda hdd    ATA_ST1000DM010-2EP1_ZN10TT3R    931G Yes        19m ago
ceph-03   /dev/sda hdd    ATA_ST1000DM010-2EP1_W9AF2W0C    931G Yes        16m ago
```

Có thể chỉ định từng ổ trên node nào triển khai OSD hoặc thêm tất cả các ổ phù hợp bằng lệnh:

```
root@ceph-01:~# ceph orch apply osd --all-available-devices
```

Đợi một lúc rồi có thể kiểm tra trạng thái

```
root@ceph-01:~# ceph osd stat
```

```
root@ceph-01:~# ceph osd tree
```

```
root@ceph-01:~# ceph orch apply osd --all-available-devices
Scheduled osd.all-available-devices update...
root@ceph-01:~# ceph osd stat
3 osds: 3 up (since 17s), 3 in (since 59s); epoch: e17
root@ceph-01:~# ceph osd tree
ID CLASS WEIGHT TYPE NAME STATUS REWEIGHT PRI-AFF
-1          2.72910 root default
-7          0.90970 host ceph-01
 2    hdd  0.90970      osd.2     up   1.00000 1.00000
-5          0.90970 host ceph-02
 1    hdd  0.90970      osd.1     up   1.00000 1.00000
-3          0.90970 host ceph-03
 0    hdd  0.90970      osd.0     up   1.00000 1.00000
root@ceph-01:~#
```

Cấu hình

```
root@ceph-01:~# ceph orch apply mgr --placement="3 label: admin"
```

```
root@ceph-01:~# ceph orch apply mon --placement="3 label: admin"
```

2 lệnh này để yêu cầu Ceph triển khai Ceph Manager(MGR) và Ceph Monitor(MON) tối đa 3 node trên các host có gắn label:_admin

Các cấu hình global khác

```
root@ceph-01:~# ceph config set global auth_client_required cephx
```

```
root@ceph-01:~# ceph config set global auth_cluster_required cephx
```

```
root@ceph-01:~# ceph config set global auth_service_required cephx
```

```
root@ceph-01:~# ceph config set global osd_pool_default_pg_autoscale_mode
off
```

```
root@ceph-01:~# ceph config set global osd_crush_update_on_start false
```

```
root@ceph-01:~# ceph orch apply mgr --placement="3 label:_admin"
Scheduled mgr update...
root@ceph-01:~# ceph orch apply mon --placement="3 label:_admin"
Scheduled mon update...
root@ceph-01:~# ceph config set global auth_client_required cephx
root@ceph-01:~# ceph config set global auth_cluster_required cephx
root@ceph-01:~# ceph config set global auth_service_required cephx
root@ceph-01:~# ceph config set global osd_pool_default_pg_autoscale_mode off
root@ceph-01:~# ceph config set global osd_crush_update_on_start false
root@ceph-01:~#
```

Kiểm tra cụm

```
root@ceph-01:~# ceph health
```

```
root@ceph-01:~# ceph health
HEALTH_OK
```

Lệnh này trả về ngắn gọn về trạng thái hiện tại của cụm Ceph

HEALTH_OK: Cụm Ceph đang hoạt động bình thường, không có vấn đề gì.

```
root@ceph-01:~# ceph -s
```

Lệnh này trả về chi tiết trạng thái của cụm Ceph, bao gồm các thông tin về health, trạng thái của các daemon (OSD, MON, MGR), và thông tin chi tiết về dung lượng lưu trữ và PGs.

```
root@ceph-01:~# ceph -s
cluster:
  id: 41073fa0-795e-11ef-bedb-fd1565f7fbc9
  health: HEALTH_OK

  services:
    mon: 3 daemons, quorum ceph-01,ceph-02,ceph-03 (age 2h)
    mgr: ceph-01.iexdoz(active, since 3h), standbys: ceph-02.cmcgcl, ceph-03.pslmxrf
    osd: 3 osds: 3 up (since 103m), 3 in (since 104m)

  data:
    pools: 1 pools, 1 pgs
    objects: 2 objects, 961 KiB
    usage: 82 MiB used, 2.7 TiB / 2.7 TiB avail
    pgs: 1 active+clean
```

```
root@ceph-01:~# ceph orch host ls
```

Lệnh này liệt kê tất cả các host trong cluster

```
root@ceph-01:~# ceph orch host ls
HOST      ADDR      LABELS  STATUS
ceph-01  192.168.3.71 _admin
ceph-02  192.168.3.72 _admin
ceph-03  192.168.3.73 _admin
3 hosts in cluster
```

```
root@ceph-01:~#ceph orch ps
```

Lệnh này cung cấp thông tin chi tiết về các tiến trình trong Ceph

```
root@ceph-01:~# ceph orch ps
NAME      HOST      PORTS      STATUS      REFRESHED      AGE      MEM USE      MEM LIM      VERSION      IMAGE ID      CONTAINER ID
alertmanager.ceph-01  ceph-01  *:9093,9094  running (104m)  76s ago      3h      13.2M      -      0.25.0      c8568f914cd2  c984de314c10
ceph-exporter.ceph-01  ceph-01      running (3h)    76s ago      3h      10.9M      -      18.2.2      3c937764e6f5  bbe150ca1c12
ceph-exporter.ceph-02  ceph-02      running (2h)    76s ago      2h      10.2M      -      18.2.2      3c937764e6f5  910121c1711f6
ceph-exporter.ceph-03  ceph-03      running (2h)    76s ago      2h      10.5M      -      18.2.2      3c937764e6f5  59b018095eb1
crash.ceph-01  ceph-01      running (3h)    76s ago      3h      7107k      -      18.2.2      3c937764e6f5  3ab635abe0f5
crash.ceph-02  ceph-02      running (2h)    76s ago      2h      7115k      -      18.2.2      3c937764e6f5  aa47a2d98697
crash.ceph-03  ceph-03      running (2h)    76s ago      2h      7128k      -      18.2.2      3c937764e6f5  1b69ef755d35
grafana.ceph-01  ceph-01  *:3000      running (3h)    76s ago      3h      72.5M      -      9.4.7       954c08fa6188  7ceb5c8a62a9
mgr.ceph-01.iexdoz  ceph-01  *:9283,8765,8443  running (3h)    76s ago      3h      497M      -      18.2.2      3c937764e6f5  78ff74e9af33
mgr.ceph-02.cmcgcl  ceph-02  *:8443,9283,8765  running (2h)    76s ago      2h      435M      -      18.2.2      3c937764e6f5  7d5b998060b3
mgr.ceph-03.pstlxf  ceph-03  *:8443,9283,8765  running (105m)  76s ago  105m      435M      -      18.2.2      3c937764e6f5  dd869ef7bea7
mon.ceph-01  ceph-01      running (3h)    76s ago      3h      94.9M      2048M     18.2.2      3c937764e6f5  9e3d1c58edd7
mon.ceph-02  ceph-02      running (2h)    76s ago      2h      72.6M      2048M     18.2.2      3c937764e6f5  1748b39e683c
mon.ceph-03  ceph-03      running (2h)    76s ago      2h      72.5M      2048M     18.2.2      3c937764e6f5  0da6a335fe13  1b4c1f534088
node-exporter.ceph-01  ceph-01  *:9100      running (3h)    76s ago      3h      9579k      -      1.5.0       0da6a335fe13  880c3a022337
node-exporter.ceph-02  ceph-02  *:9100      running (2h)    76s ago      2h      9287k      -      1.5.0       0da6a335fe13  40e4e0d8e8cf
node-exporter.ceph-03  ceph-03  *:9100      running (2h)    76s ago      2h      9367k      -      1.5.0       0da6a335fe13  ca891c1c27ac
osd.0   ceph-03      running (115m)  76s ago  115m      64.1M      3775M     18.2.2      3c937764e6f5
osd.1   ceph-02      running (115m)  76s ago  115m      64.0M      3775M     18.2.2      3c937764e6f5
osd.2   ceph-01      running (115m)  76s ago  115m      64.2M      4096M     18.2.2      3c937764e6f5
prometheus.ceph-01  ceph-01  *:9095      running (2h)    76s ago      3h      58.4M      -      2.43.0      a07b618ecd1d  8bda4e799f23
```

```
root@ceph-01:~#ceph config dump
```

Lệnh này trả về cấu hình của cụm

```
root@ceph-01:~# ceph config dump
WHO      MASK      LEVEL      OPTION      VALUE      RO
global      advanced auth_client_required      cephx
global      advanced auth_cluster_required      cephx
global      advanced auth_service_required      cephx
global      advanced cluster_network      10.10.10.0/24
global      basic container_image      quay.io/ceph/ceph@sha256:f8d467dcf49d13b8ea42229d89be642581110175d8ce36e216aefc9b32b0854d
global      advanced osd_crush_update_on_start      false
global      advanced osd_pool_default_pg_autoscale_mode      off
global      advanced auth_allow_insecure_global_id_reclaim      false
mon      advanced public_network      192.168.3.0/24
mgr      advanced mgr/cephadm/container_init      True
mgr      advanced mgr/cephadm/migration_current      6
mgr      advanced mgr/dashboard/ALERTMANAGER_API_HOST      http://ceph-01:9093
mgr      advanced mgr/dashboard/GRAFANA_API_SSL_VERIFY      false
mgr      advanced mgr/dashboard/GRAFANA_API_URL      https://ceph-01:3000
mgr      advanced mgr/dashboard/PROMETHEUS_API_HOST      http://ceph-01:9095
mgr      advanced mgr/dashboard/ssl_server_port      8443
mgr      advanced mgr/orchestrator/orchestrator      cephadm
osd      host:ceph-02 basic osd_memory_target      3958753280
osd      host:ceph-03 basic osd_memory_target      3958678732
osd      advanced osd_memory_target_autotune      true
osd.0   basic osd_mclock_max_capacity_iops_hdd      225.563067
osd.1   basic osd_mclock_max_capacity_iops_hdd      227.717365
osd.2   basic osd_mclock_max_capacity_iops_hdd      212.654862
```